

Title: Implementing Software Timers  
 By: Don Libes

Originally appeared in the Nov. 1990 "C User's Journal" and is also reprinted as Chapter 35 of "Obfuscated C and Other Mysteries", John Wiley & Sons, 1993, ISBN 0-471-57805-3.  
<http://www.wiley.com/compbooks/m3.html>.

---

This column will describe a set of functions to implement software timers. What are software timers and why do you need them? Software timers make up for inherent limitations in hardware timers. For example, while most computers have clock hardware, you can typically only have the clock trigger an interrupt for one time in the future.

When running multiple tasks, you will want to have the clock keep track of multiple timers concurrently so that interrupts can be generated correctly even if the time periods overlap. Operating systems do this all the time.

Robert Ward discussed the related problem of building a general purpose scheduler ("Practical Schedulers for Real-Time Applications") in the April 1990 CUJ. In the "Additional Ideas" section, Robert described the usefulness of a timer scheduling queue. "Events can specify the timing of other events by putting a timer programming request in a special queue." That is exactly what the software in this column will do. (Thanks for the lead in, Robert.) You may want to reread at least the beginning of his column right now, although it isn't really necessary.

The code in this column has other uses as well. For example, you can use it to simulate multiple timers in environments such as a UNIX process which only allows the user one software timer. Even if you aren't interested in software timers, I think you will find this an intriguing column. Using simple techniques and data structures, this C code produces very powerful results. The code was very tricky to get right, and my commentary should be interesting if only as some more practice in reading and writing C code.

## Timers

By implementing the timers as a separate piece of software, we can reduce the complexity of the scheduler. Some people like this kind of modularization, and some don't. Similarly some operating systems do this, and some don't. I like it. It makes the code easier to write, to read, and to correct (oops).

The basic idea of a timer is that they allow tasks to be run at some time in the future. When their time arrives, they are scheduled to be run. The responsibility of actually running them is then turned over to someone else, such as the scheduler. In order to communicate with the scheduler, we'll set up a common data structure called a timer (listing 1). I've also included a few other miscellaneous definitions that will be needed later on. For instance, the TIME typedef is used to declare all relative time variables. You can complete this definition based on what your needs are.

```
#include <stdio.h>

#define TRUE    1
#define FALSE   0

#define MAX_TIMERS    ...    /* number of timers */
typedef ... TIME;      /* how time is actually stored */
#define VERY_LONG_TIME ...    /* longest time possible */

struct timer {
    int inuse;          /* TRUE if in use */
    TIME time;          /* relative time to wait */
```

```

    char *event;          /* set to TRUE at timeout */
} timers[MAX_TIMERS];    /* set of timers */

```

#### listing 1

Each timer will be represented by a timer struct. The set of timers will be maintained in an array, `timers`. The first element of each timer declares whether the timer is in use. The second element of a timer is the amount of time being waited for. As time passes, this will be periodically updated. `event` is a pointer to a value that is initially set to 0. When it is time to run the task, `*event` is set to 1. We can imagine that the scheduler also keeps an event pointer. Every so often, it reexamines it. When it finds it has been set to 1, it knows that the timer has expired and the associated task can be run.

[Notice how simple this is. Other schedulers or other scheduler data structures could enable runnability, without worrying or even knowing about timers.]

The code in listing 2 initializes the timers. It runs through the array setting each `inuse` flag to `FALSE`. This for loop will become idiomatic to you by the end of this column.

```

void
timers_init() {
    struct timer *t;

    for (t=timers;t<&timers[MAX_TIMERS];t++)
        t->inuse = FALSE;
}

```

#### listing 2

Now we can write the routines to schedule the timers. First, I'll show `timer_undeclare`, which is a little simpler than its counterpart, `timer_declare`.

There are a variety of ways to keep track of the timers. Machines which don't have sophisticated clock hardware usually call an interrupt handler at every clock tick. The software then maintains the system time in a register, as well as checking for timer entries that have expired.

More intelligent machines can maintain the clock in hardware, only interrupting the CPU after a given time period has expired. By having the clock interrupt for when an event is waiting, you can get a tremendous speedup. This technique is also common in software simulations and thread implementation.

Reading the clock may require an operating system call, but for our purposes we will assume the variable `time_now` to be automatically updated by the hardware for just this purpose. `volatile` indicates that the variable should not be cached in a register but read from storage each time.

```
volatile TIME time_now;
```

We will define several variables for shorthands. `timer_next` will point to the timer entry that we next expect to expire. `time_timer_set` will contain the system time when the hardware timer was last set.

```

struct timer *timer_next = NULL; /* timer we expect to run down next */
TIME time_timer_set;           /* time when physical timer was set */

void timers_update();          /* see discussion below */

void

```

```

timer_undecare(t)
struct timer *t;
{
    disable_interrupts();
    if (!t->inuse) {
        enable_interrupts();
        return;
    }

    t->inuse = FALSE;

    /* check if we were waiting on this one */
    if (t == timer_next) {
        timers_update(time_now - time_timer_set);
        if (timer_next) {
            start_physical_timer(timer_next->time);
            time_timer_set = time_now;
        }
    }
    enable_interrupts();
}

```

### Listing 3

#### Undeclaring Timers - Why and How?

`timer_undecare` does just what its name implies, it undeclares a timer. Undeclaring timers is actually an important operation in some applications. For example, network code sets timers like crazy. In some protocols, each packet sent generates a timer. If the sender doesn't receive an acknowledgement after a given interval, the timer forces it to resend a packet. If the sender does receive an acknowledgement, it undeclares the timer. If things are going well, every single timer declared is later undeclared.

`timer_undecare` (listing 3) is performed with interrupts disabled. This is necessary because we are going to have an interrupt handler that can access the same data. Because this data is shared, access must be strictly controlled. I've shown the interrupt manipulation as a function call, but you must use whatever is appropriate to your system. This is very system dependent.

`timer_undecare` starts by checking the validity of the argument as a timer entry. We will see later that the system clock can implicitly undeclare timer entries. Thus we must make a reasonable attempt to assure ourselves that a timer to be undeclared is still declared.

Once assured the timer is valid, `timer_undecare` marks the entry invalid. If the timer happens to be the very one next expected to expire, the physical timer must be restarted for the next shorter timer. Before doing that, all the timer entries have to be updated by the amount of time that has elapsed since the timer was last set. This is done by `timers_update` which also calculates the next shortest timer. Looking for the shortest timer in that function is a little obscure but happens to be very convenient since `timers_update` has to look at every timer anyway.

`timers_update` (listing 4) goes through the timers, subtracting the given time from each. If any reach 0 this way, they are triggered by setting the event flag. Any lag in the difference between when a timer was requested and `timers_update` is called, is accounted for by basing the latency against `time_now` and also collecting timers that have "gone negative" in `timers_update`. (Why might a timer go negative?) Lastly, we also remember the lowest nonzero timer to wait for as `timer_next`.

`timer_last` is just a temporary. It is a permanently non-schedulable timer that will only show up when all the other timers have been scheduled.

```

/* subtract time from all timers, enabling any that run out along the way */
void
timers_update(time)
TIME time;
{
    static struct timer timer_last = {
        FALSE                /* in use */,
        VERY_LONG_TIME       /* time */,
        NULL                 /* event pointer */
    };

    struct timer *t;

    timer_next = &timer_last;

    for (t=timers;t<&timers[MAX_TIMERS];t++) {
        if (t->inuse) {
            if (time < t->time) { /* unexpired */
                t->time -= time;
                if (t->time < timer_next->time)
                    timer_next = t;
            } else { /* expired */
                /* tell scheduler */
                *t->event = TRUE;
                t->inuse = 0; /* remove timer */
            }
        }
    }

    /* reset timer_next if no timers found */
    if (!timer_next->inuse) timer_next = 0;
}

```

listing 4

#### Declaring Timers

timer\_declare (listing 5) takes a time and an event address as arguments. When the time expires, the value that event points to will be set. (This occurs in timers\_update under the comment /\* tell scheduler \*/.) timer\_declare returns a pointer to a timer. This pointer is the same one that timer\_undecare takes as an argument.

```

struct timer *
timer_declare(time,event)
unsigned int time;          /* time to wait in 10msec ticks */
char *event;
{
    struct timer *t;

    disable_interrupts();

    for (t=timers;t<&timers[MAX_TIMERS];t++) {
        if (!t->inuse) break;
    }

    /* out of timers? */
    if (t == &timers[MAX_TIMERS]) {
        enable_interrupts();
        return(0);
    }

    /* install new timer */
    t->event = event;
    t->time = time;
    if (!timer_next) {
        /* no timers set at all, so this is shortest */
        time_timer_set = time_now;
    }
}

```

```

        start_physical_timer((timer_next = t)->time);
    } else if ((time + time_now) < (timer_next->time + time_timer_set)) {
        /* new timer is shorter than current one, so */
        timers_update(time_now - time_timer_set);
        time_timer_set = time_now;
        start_physical_timer((timer_next = t)->time);
    } else {
        /* new timer is longer, than current one */
    }
    t->inuse = TRUE;
    enable_interrupts();
    return(t);
}

```

#### listing 5

As with its counterpart, interrupts are disabled in timer\_declare to prevent concurrent access to the shared data structure.

The first thing timer\_declare does is to allocate a timer. If none are available, a NULL is returned so that the caller can fail or retry later.

Once a timer is allocated and initialized, we must check if the physical timer must be changed. There are three cases:

- 1) There are no other timers;

In this case, we go ahead and start the physical timer with the time of this timer.

- 2) There are other timers, but this new one is the shortest of all the others;

In this case, we must restart the physical timer to the new time. But before we do that, we must update all the other timers by the amount of time that has elapsed since the physical timer was last set.

- 3) There are other timers, and this new one is not the shortest.

There is nothing to do in this case. However, for legibility it is broken into its own case which contains only a comment. That way it is clear what is going on when the previous else-if test fails.

Before enabling interrupts and returning, the timer's inuse flag is set. The reason it is done afterwards rather than with the earlier timer settings is that this prevents timers\_update from updating it with a time period that occurred before it was even declared.

#### Handling Timer Interrupts

The only remaining routine is the interrupt handler (listing 6) actually called when the physical clock expires. When the interrupt handler is called, we are guaranteed that the time described by timer\_next has elapsed.

```

void
timer_interrupt_handler() {
    timers_update(time_now - time_timer_set);

    /* start physical timer for next shortest time if one exists */
    if (timer_next) {
        time_timer_set = time_now;
        start_physical_timer(timer_next->time);
    }
}

```

#### listing 6

Each time the interrupt handler is called, a timer has expired. By calling `timers_update`, all the timers will be decremented and any timers that have expired will have their event flags enabled. This will also set up `timer_next` so that the physical timer can be restarted for the next timer we expect to occur.

Let's examine one special case. Suppose we have only one timer set up. Now imagine that we have called `timer_undeclare` and just as interrupts are disabled, the physical clock ticks down all the way. Since interrupts are disabled, the interrupt will be delivered immediately after interrupts are enabled. But they will be enabled after the timer has been deleted. So we see a situation where an interrupt will be delivered for a timer that no longer exists. What occurs in the interrupt handler?

`timers_update` is called. It finds nothing to update. As a consequence of this, `timer_next` is set to 0. The remainder of the interrupt handler already handles the case of no remaining timers, and the handler returns normally.

This is an example of the kind of special casing you have to keep in mind when writing the code. (In fact, my first implementation didn't handle this right, and it was painful to debug. Debuggers don't work very well when fooling around with interrupts!)

## Conclusion

I have presented an implementation of timers. The code is carefully designed so that it is relatively free of special demands it places on a scheduler. For example, it doesn't close off the scheduler from using a different kind of timer at the same time.

One thought that may have occurred to you while reading this, is why the timers are maintained as an array rather than say, a linked list. Using a linked list would avoid the overhead of stepping through arrays (which can be almost entirely empty). Keeping the list sorted by time would make the `timers_update` function much simpler.

On the other hand, it would complicate the other functions. For example, `timer_undeclare`, would either require you to use a doubly-linked list, or to search the entire list from the beginning each time. Another point is that real-time systems typically avoid dynamic structures to begin with. For example, using `malloc/free` from a process-wide heap can cause an indeterminate amount of time that is difficult to estimate. If I was to recode this using linked lists, I would use a `malloc` implementation from a small pool of timer-only buffers, which in effect is very similar to what I've done here with arrays. There would be a tradeoff in space and time, which you might prefer or not depending upon your application.

If you decide to recode or just modify my implementation, be very careful. Always imagine the worst thing that can happen when two processes attempt to access the same data structure at the same time. Happy interruptions!

Thanks

Debugging timing routines is very different than other code, since unrelated events in the computer can make your programs behave differently. Even putting in `printf` statements can change critical execution paths. It is extremely aggravating when problems disappear only when you are debugging. Furthermore, most debuggers do not work well when interrupts are disabled. Ed Barkmeyer was of great help debugging the timer code and teaching me to persevere when I saw code behaving in ways that had to be impossible. Thanks to Sarah Wallace and Randy Miller who debugged this column and also forced me to make all the explanations much clearer.